SENTINEL-X: WORKING & TESTING GUIDE

Defensive Attack Forecasting SOC • End-to-End Pipeline & Automated Verification

Project: Sentinel-X (SIH26153)	**Department:** Computer Science & Engineering
Institution: Buddha Institute of Technology, Gorakhpur (AKTU)	**Guide:** Mr. Akash Gupta (Assistant Professor)
Team: Shristi Jaiswal (2405250100153) • Suraj Prakash Chaudhary (2405250100156)	Suraj Upadhyay (2405250100158) • Udai Kumar Srivastava (2405250100160)
Live Dashboard URL: http://localhost:8501	**Repository:** github.com/surajbuilds53/sentinel-x

1. Executive Summary & The Predictive Mental Model

The Fundamental Problem: Traditional intrusion detection systems (NIDS like Snort and Suricata) are **purely reactive**. An alarm sounds only *after* a packet payload has executed and compromised the target machine. According to industry benchmarks, the average Mean Time to Detect (MTTD) is **207 days**, and SOC analysts face severe alert fatigue from thousands of daily false positives.

The Sentinel-X Paradigm Shift: Sentinel-X acts as a **predictive weather radar for cyber threats**. Rather than evaluating isolated static packets, Sentinel-X tracks continuous multi-dimensional network kinetics over rolling 120-second historical sequences. By learning the physical dynamics of cyberattacks, it **forecasts attack probability 10 to 30 seconds into the future (horizons T+1, T+2, T+3)**. This anticipatory lead-time enables automated firewall Access Control Lists (ACLs) to isolate malicious IPs *before* a breach succeeds.

Pillar	Traditional Reactive NIDS	Sentinel-X Predictive Defense
Detection Timing	Post-compromise (fires after exploit executes)	**Pre-emptive (forecasts 10–30s ahead of exploit)**
Input Telemetry	Isolated, static packet headers / signatures	**16 continuous temporal features (rolling 120s window)**
Model Architecture	Static rules, Decision Trees, or 1-step classifiers	**2-Layer LSTM World Model + Autoregressive Rollout**
False Alarms	High false positive rate (~20%), causing alert fatigue	**75% reduction in FPR (down to 5.0%) with 100% T+1 recall**
Explainability	Opaque black-box score or raw packet dump	**Dual-tier SHAP feature attribution + temporal saliency**
Edge Execution	Heavy cloud compute dependencies & latency	**Sub-20ms local CPU execution; operates 100% offline**

2. Step-by-Step Architecture: The 6 Pipeline Stages

Stage 1: 10-Second Feature Aggregation (sentinel_x/data/aggregation.py)

Individual raw packets are too noisy to reveal macro behavioral intent. Sentinel-X aggregates continuous flow telemetry into **10-second non-overlapping time buckets**. For each window, it computes **16 mathematical state features** capturing volume kinetics, packet duration distributions, inter-arrival intervals (which expose robotic bot loops), and TCP flag distributions (SYN, ACK, RST, FIN, PSH ratios).

Stage 2: Chronological Sequence Windowing (sentinel_x/data/windowing.py)

To give the model temporal memory, Sentinel-X groups twelve consecutive 10-second windows into a **120-second rolling historical sequence** (time-steps T-11 through T), forming a tensor of dimension **[12, 16]**. Crucially, dataset partitioning follows a strict **chronological 70/15/15 split** (first 70% for training, next 15% for validation, last 15% for testing). Random shuffling is forbidden, preventing future data leakage and ensuring realistic generalization to unseen attacks.

Stage 3: The 2-Layer LSTM World Model (sentinel_x/models/lstm_world_model.py)

The neural backbone is a **2-layer Long Short-Term Memory (LSTM) recurrent network** with 128 hidden units and 0.3 dropout. The network features three multi-task heads:

• **Risk Head (Binary Cross-Entropy):** Forecasts probability of intrusion $\in$ [0, 1].
• **Stage Head (Categorical Cross-Entropy):** Classifies traffic into 5 MITRE ATT&CK; stages.
• **Reconstruction Head (Mean Squared Error):** Reconstructs the 16-D feature vector of the next time step, forcing the model to learn the true physical dynamics of network behavior.

Stage 4: Multi-Step Autoregressive Rollout (Iterative Horizon Projection)

To forecast beyond immediate 10-second horizons, Sentinel-X applies **autoregressive state rollout**:
1. At T+1 (+10s), the model predicts risk and reconstructs the future state vector x_{T+1}.
2. It appends x_{T+1} to the sequence, discards the oldest time step T-11, and feeds the updated 120s sequence back into the LSTM to forecast T+2 (+20s).
3. It repeats this process to forecast T+3 (+30s), enabling proactive multi-step defense.

Stage 5: Dual-Tier Explainable AI with SHAP (sentinel_x/explainability/shap_explainer.py)

To eliminate the dangerous 'black box' problem, Sentinel-X integrates game-theoretic **Shapley Additive exPlanations (SHAP)**:
• **Signed Threat Drivers (Red Bars):** Identifies exact features pushing risk upward (e.g., sudden syn_ratio or packet_rate surges).
• **Threat Inhibitors (Green Bars):** Highlights stabilizing features pulling risk downward (e.g., high ack_ratio, normal flow durations).
• **Temporal Sequence Saliency:** Mathematically proves which historical time-steps in the 120s sequence triggered the model's decision.

Stage 6: Interactive SOC Command Dashboard (dashboard/app.py)

A production-grade web application built in **Streamlit and Plotly**, accessible locally at **http://localhost:8501**. Features 5 specialized operational views, sub-20ms inference latency, real-time scenario simulation, and automated mitigation advice.

3. The 16 Continuous State Features & Cyber Significance

Feature Name	Category	Measurement / Formula	Cyber Defense Significance
flow_count	Volume	Total flows in 10s window	Spikes indicate port scanning or DDoS amplification.
tot_fwd_bytes	Volume	Sum of client → server bytes	Surges reveal payload delivery or brute-force requests.
tot_bwd_bytes	Volume	Sum of server → client bytes	Elevated volumes expose unauthorized data exfiltration.
packet_rate	Velocity	Packets per second	Bursts expose high-frequency automated attack tools.
byte_rate	Velocity	Bytes per second	Abnormal throughput flags high-speed network abuse.
duration_mean	Dynamics	Average flow duration (sec)	Extremely short flows indicate scanning; long flows indicate C2.
duration_std	Dynamics	Standard deviation of duration	Variance highlights heterogeneous background traffic.
fwd_pkt_len_mean	Dynamics	Mean forward packet size	Abnormally small/large sizes expose exploit framing.
bwd_pkt_len_mean	Dynamics	Mean backward packet size	Captures server reply dynamics and exfiltration payload size.
iat_mean	Timing	Mean packet inter-arrival time	Very low values flag rapid volumetric flooding.
iat_std	Timing	Variance of inter-arrival times	Near-zero variance exposes periodic robotic bot beacons.
syn_ratio	TCP Flags	SYN packets / Total packets	Dominant ratio indicates SYN floods and TCP port probes.
ack_ratio	TCP Flags	ACK packets / Total packets	High ratio confirms established, benign two-way communication.
rst_ratio	TCP Flags	RST packets / Total packets	Spikes expose rejected ports and terminated connection attempts.
fin_ratio	TCP Flags	FIN packets / Total packets	Monitors orderly TCP connection teardowns.
psh_ratio	TCP Flags	PSH packets / Total packets	Elevated levels indicate urgent push payloads and interactive shells.

4. Step-by-Step Testing & Verification Guide

A. Interactive Testing via Browser Dashboard (http://localhost:8501)

The primary interface for presentation and evaluation is the **SOC Command Dashboard**. It runs locally and connects directly to the trained PyTorch model.

Step 1: Launch and Verify Server

If the server is not already active, launch it from the project root using:

```
.\.venv\Scripts\python.exe -m streamlit run dashboard/app.py --server.port 8501
```

Open your browser and navigate to **http://localhost:8501**. The application should load instantly without external internet dependencies.

Step 2: Test Scenario 1 — Normal Benign Baseline

In the left sidebar, ensure **'Enable Demo Mode'** is checked. Select **'Normal Benign Baseline'**.
• *Observe View 1 (Executive Risk):* All three risk gauges (T+1, T+2, T+3) remain **green (< 40%)**.
• *Stage Badge:* Displays **'BENIGN (Stage 0)'** in green.
• *Automated Mitigation:* Displays 'Normal network operations. Zero defensive intervention required.'

Step 3: Test Scenario 2 — Reconnaissance (PortScan / Probing)

Select **'Reconnaissance (PortScan / Probing)'** from the scenario dropdown.
• *Observe View 1:* The T+1 gauge swings immediately to **elevated risk** (> 75%, Red).
• *Stage Badge:* Automatically switches to **'RECONNAISSANCE (Stage 1)'** with high confidence.
• *Automated Mitigation:* Displays 'Pre-emptively apply firewall ACL rule to isolate source IP.'

Step 4: Test Scenario 3 — Initial Access (Brute Force / Patator)

Select **'Initial Access (Patator / Breach)'**.
• *Observe View 1:* Demonstrates aggressive escalation as repeated authentication bursts occur.
• *Stage Badge:* Changes to **'INITIAL ACCESS (Stage 2)'**.

Step 5: Test View 2 — Forecast Timeline

Select **'Forecast Timeline (View 2)'** in the sidebar.
• *Cyan Line:* Shows the 120-second historical sequence of observed network risk (T-11 to T).
• *Dashed Red Line:* Shows the 3-step autoregressive future forecast (T+1, T+2, T+3).
• *Shaded Band:* Represents model prediction confidence intervals, visualising variance over expanding forecast horizons.

Step 6: Test View 3 — Threat Attribution & Explainable AI (SHAP)

Select **'Threat Attribution (View 3)'** in the sidebar.
• *Horizontal Bar Chart:* Displays signed SHAP feature attributions. Red bars (positive impact) expose specific attack drivers (e.g., syn_ratio and packet_rate spikes), while green bars expose inhibitors.
• *Temporal Saliency:* Bar chart confirms that recent kinetic changes (T-1 and T) carry the highest attribution weights.

Step 7: Test View 4 & 5 — Model Performance & System Diagnostics

• *View 4 (Model Performance):* Confirms the 100% recall at T+1 and 75% FPR reduction.
• *View 5 (System Status):* Displays hardware compute status, memory consumption, and verifies edge execution capability.

B. Automated Code Testing via Terminal (CLI)

To verify all mathematical formulations, data aggregations, and neural weights independently of the browser UI, execute the following commands in PowerShell from the project root:

1. Run Complete Automated Pytest Suite

```
.\.venv\Scripts\python.exe -m pytest tests -v
```

Executes 29 unit and integration tests across data processing, windowing, LSTM forward passes, SHAP attributions, and dashboard rendering. **Status: 29 Passed (100% Pass Rate).**

2. Run Comparative Baseline Evaluation

```
.\.venv\Scripts\python.exe scripts/evaluate.py
```

Evaluates both Logistic Regression baseline and the 2-Layer LSTM World Model on held-out test data. Generates precision, recall, F1-score, and confusion matrices.

3. Audit System Environment & Dependencies

```
.\.venv\Scripts\python.exe scripts/check_environment.py
```

Verifies installation of PyTorch, NumPy, Pandas, Scikit-Learn, SHAP, Streamlit, Plotly, ReportLab, and confirms all project directory paths.

4. Re-run End-to-End Data Preparation

```
.\.venv\Scripts\python.exe scripts/prepare_data.py
```

Regenerates the 10s kinetic aggregations and 70/15/15 chronological sequence tensors into data/processed/processed_windows.npz.

5. Empirical Benchmarks & Viva Defense Summary

Evaluation Metric	Baseline (Logistic Regression)	Sentinel-X (LSTM World Model)	Defensive Advantage
Recall at T+1 (+10s)	97.8% (missed 2.2% of attacks)	**100.0% (Zero missed attacks)**	Zero false negatives; critical intrusions never slip through.
False Positive Rate (FPR)	20.0% (1 in 5 normal flows flagged)	**5.0% (75% relative reduction)**	Solves SOC analyst alert fatigue and avoids costly false blocks.
F1-Score at T+1	0.8950	**0.9834**	+8.84% overall performance improvement over standard baselines.
F1-Score at T+2 (+20s)	N/A (cannot forecast multi-step)	**0.9765**	Maintains strong predictive stability across expanding horizon.
F1-Score at T+3 (+30s)	N/A (cannot forecast multi-step)	**0.9721**	Graceful degradation; reliable pre-emptive lead time.
Inference Latency	2.1 ms	**18.4 ms (Standard CPU)**	Sub-20ms requirement met without needing cloud GPU servers.
Automated Test Suite	N/A	**29 Passed / 0 Failed (100%)**	Fully verified across all 9 unit and integration modules.

Key Viva Examination Defense Points:

Q: Why an LSTM rather than Random Forest or XGBoost?

Answer: Cyberattacks have a physical arrow of time. A port scan precedes a brute force attempt, which precedes lateral movement. Static tree models look only at isolated moments and cannot model continuous temporal kinetics. An LSTM maintains internal cell states that preserve 120-second temporal dependencies and enable autoregressive future rollout.

Q: What is data leakage and how was it avoided?

Answer: Randomly shuffling time-series data leaks future packet structures into the training set, causing artificially inflated metrics. Sentinel-X enforces a strict chronological 70/15/15 time partition, proving that our 100% recall is authentic and evaluated on unseen future timestamps.

Q: How does SHAP benefit SOC analysts?

Answer: If a firewall abruptly blocks an IP, analysts must know why. SHAP decomposes the neural output into signed numerical contributions. If the top driver is syn_ratio at +0.35 alongside packet_rate at +0.25, the analyst can confidently confirm a SYN-flood scan and approve the block rule.